

Lecture 8 - Applications of Image Processing in Machine Learning

January 31, 2021

1 Lecture 8 - Applications of Image Processing in Machine Learning

In this notebook, we will apply image processing in preparing our dataset for a machine learning algorithm.

Authored by: Benjur Emmanuel L. Borja

1.1 Leaf Classification

In the previous lectures, we discussed how to visualize, enhance, and extract information from images. We now apply the different algorithms that we have learned to create a leaf classification algorithm. Let's try it out. The file *leaves.zip* contains images of five classes of leaves. We need to create a pipeline that would 1. Read and clean the images 1. Segment objects of interest 1. Extract features from the objects 1. Train an ML model

Take note that each of the steps can affect the accuracy of our model. Prioritizing which step to update will rely on heuristics. Let's practice:

Breakout Room Exercise: Create a machine learning pipeline and try to reach the highest accuracy possible. Make sure that you are able to answer the corresponding questions. 1. Did you apply any data cleaning methods to prepare the images? If yes, what are these methods? Upon inspection of the images, it was observed that there are some artifacts from the scanning process. To address this, we first convert the image to grayscale, then simply binarize the result using a threshold of 0.39.

2. What algorithms did you use to segment the different leaves? What worked best? As mentioned above, there were observed artifacts from scanning, and as well as some leaf samples are touching, which results in conjoined continuous elements. Therefore, morphological operations (i.e. erosion) was used to ensure that each leaf image segments neatly, as much as possible. To address any samples that are still conjoined, we simply filter the dataset to only include up to the 95th percentile in terms of segment area.
3. What are the features that you have extracted from the leaves? What are the derived features that you have extracted? And what are the significant features in classifying each leaf? The following features were extracted using skimage's regionprops function:
 - perimeter - approximation of the image contour length
 - area - number of pixels in the region

- `bbox_area` - number of pixels contained by the bounding box
- `convex_area` - number of pixels in the convex hull image,
- `eccentricity` - ratio of the focal distance of the circumscribed ellipse to the major axis length
- `equivalent_diameter` - diameter of the circle with the same area as the region
- `extent` - ratio of area to bounding box area
- `filled_area` - number of pixels in the area, plus pixels of all the holes
- `major_axis_length` - major axis of the circumscribed ellipse
- `minor_axis_length` - minor axis of the circumscribed ellipse
- `local_centroid_y` - bounding box centroid - y
- `local_centroid_x` - bounding box centroid - x
- `euler_number` - number of connected components subtracted by number of holes
- `perimeter_crofton` - perimeter from the Crofton formula
- `solidity` - ratio of area to convex hull area

In terms of derived features, we experimented with including the ratio of perimeter to area, and the ratio of the minor axis length to the major axis length.

4. How does the limited number of samples affect your training? On the one hand, it is reasonable to say that the limited number of samples makes it easier for the machine to learn this particular dataset. However, at the same time, it is hampered in its generalizability.
5. What machine learning algorithm did you use? What is your best performing algorithm? The chosen model for the experiment is the Random Forest Classifier. This is primarily due to its speed. While the best result is when we add the **ratio between the axis lengths** (90.6), it appears that neither of the additional features improved the model significantly. The resulting test accuracies are within two percent (2%) of each other. More research and experiments are needed.

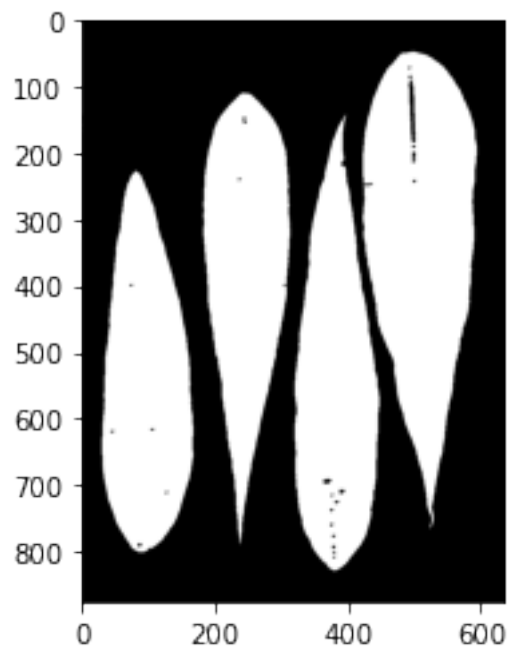
```
[75]: import os
import numpy as np
from skimage import data, io, filters
from skimage.measure import label, regionprops
from skimage.morphology import erosion, closing
from skimage.color import label2rgb, rgb2gray
import matplotlib.pyplot as plt
```

```
[76]: kernel = np.array([[0,0,1,1,1],
                        [0,1,1,1,1],
                        [1,1,1,1,1]])
```

```
[104]: leaves = {'A':[], 'B':[], 'C':[], 'D':[], 'E':[]}
for img in os.listdir(os.getcwd()):
    if img.startswith('plant'):
        leaves[img[5]].append(erosion(rgb2gray(io.imread(img)) < 0.39,
                                    kernel))
```


```
[105]: plt.imshow(leaves['A'][0], 'gray')
```

```
[105]: <matplotlib.image.AxesImage at 0x7fe6898d9a30>
```







```
[106]: sample = leaves['A'][0]
sample_labeled = label(sample)
plt.imshow(sample_labeled)
```



```
[106]: <matplotlib.image.AxesImage at 0x7fe68d2668b0>
```






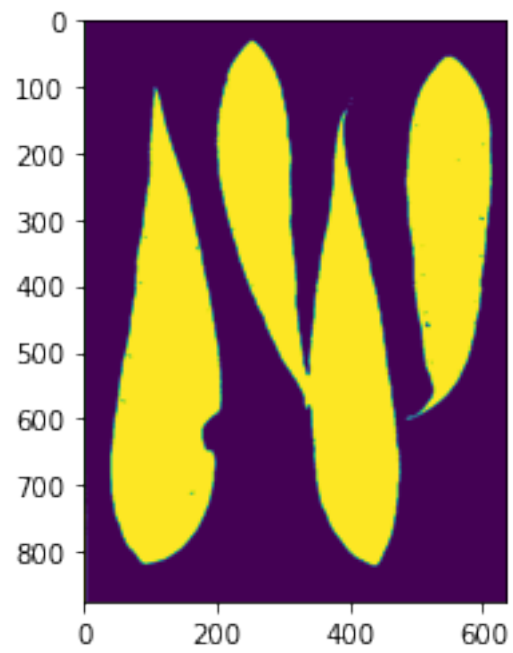
```
[152]: for img in os.listdir(os.getcwd()):  
        if img.startswith('plantA'):  
            plt.imshow(rgb2gray(io.imread(img)) < 0.39)  
            plt.show()
```



















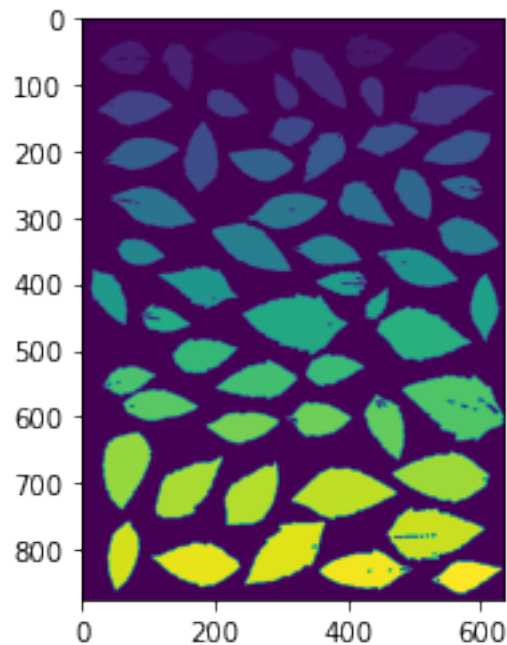
```
[109]: sample2 = leaves['B'][0]
sample2_labeled = label(sample2)
plt.imshow(sample2_labeled)
```

```
[109]: <matplotlib.image.AxesImage at 0x7fe68a8784f0>
```



```
[110]: sample3 = leaves['C'][0]
sample3_labeled = label(sample3)
plt.imshow(sample3_labeled)
```

```
[110]: <matplotlib.image.AxesImage at 0x7fe68a5ef1f0>
```




```
[111]: sample3_props=regionprops(sample3_labeled)
sample3_props=sorted(sample3_props, key=lambda x: x.area, reverse=True)[1:]
print(sample3_props[-5].area, sample3_props[-4].area)
```

3 2


```
[84]: sample3 = leaves['D'][0]
sample3_labeled = label(sample3)
plt.imshow(sample3_labeled)
```

```
[84]: <matplotlib.image.AxesImage at 0x7fe68a544bb0>
```






```
[112]: sample3 = leaves['E'][0]
sample3_labeled = label(sample3)
plt.imshow(sample3_labeled)
```



```
[112]: <matplotlib.image.AxesImage at 0x7fe689f0ffd0>
```







```
[120]: for leaf_class, images in leaves.items():
        for img in images:
            temp = label(img)
            temp_props = regionprops(temp)
            plt.imshow(temp==0)
            plt.title(f'{temp_props[0].bbox_area}')
            plt.show()
```

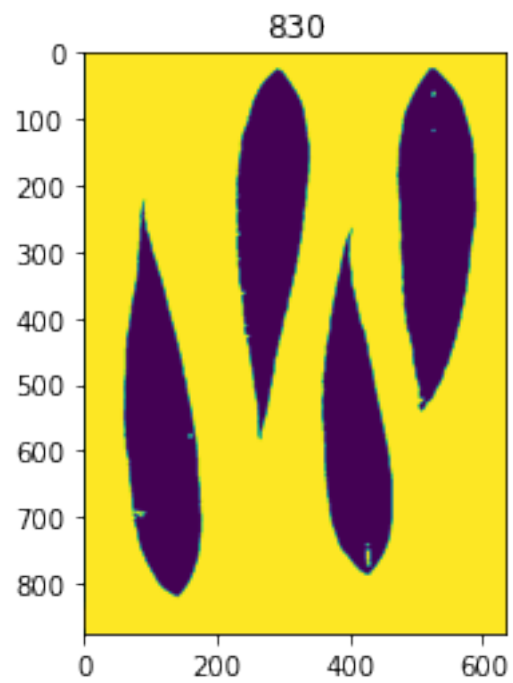













































```
[185]: # major_axis_length, local_centroid, filled_area, extent, equivalent_diameter,
# eccentricity, convex_area, bbox_area, area, perimeter, minor_axis_length
```

```

import pandas as pd

data = pd.DataFrame()

for leaf_class, images in leaves.items():
    for img in images:
        temp = label(img)
        temp_props = regionprops(temp)
        temp_props = [x for x in temp_props if x.area > 100]
        temp_props = [[x.perimeter, x.area, x.bbox_area, x.convex_area,
                        x.eccentricity, x.equivalent_diameter, x.extent,
                        x.filled_area, x.major_axis_length,
                        x.minor_axis_length, x.local_centroid[0],
                        x.local_centroid[1], x.euler_number, x.solidity,
                        leaf_class]
                       for x in temp_props]
        data = pd.concat([data, pd.DataFrame(np.array(temp_props))])

```

```

[186]: data.columns = ['perimeter', 'area', 'bbox_area', 'convex_area',
                      'eccentricity', 'equivalent_diameter', 'extent',
                      'filled_area', 'major_axis_length', 'minor_axis_length',
                      'local_centroid_y', 'local_centroid_x', 'euler', 'solidity',
                      'class']

```

```

[187]: data = data[data.area.astype(float) < data.area.astype(float).quantile(0.95)]

X = data.iloc[:, :-1].astype(float)
Y = data['class']

X.describe()

```

```

[187]:

```

	perimeter	area	bbox_area	convex_area	eccentricity \
count	283.000000	283.000000	283.000000	283.000000	283.000000
mean	663.834682	15967.310954	27985.696113	18772.547703	0.829478
std	544.688045	20940.770619	47416.278346	28204.562513	0.177992
min	45.627417	103.000000	103.000000	103.000000	0.160207
25%	262.814755	2489.000000	4183.000000	2755.000000	0.825015
50%	343.806133	3994.000000	6540.000000	4266.000000	0.865297
75%	1083.043289	31798.000000	49086.500000	35133.500000	0.972688
max	2873.572727	68600.000000	502824.000000	252726.000000	1.000000

	equivalent_diameter	extent	filled_area	major_axis_length \
count	283.000000	283.000000	283.000000	283.000000
mean	113.265694	0.629944	16019.257951	265.246759
std	86.762285	0.147923	20992.497875	280.881065
min	11.451798	0.004938	103.000000	13.565450

25%	56.294659	0.574495	2528.500000	88.501715
50%	71.311421	0.629047	3997.000000	113.795177
75%	201.205981	0.712019	32007.000000	370.029779
max	295.540577	1.000000	68686.000000	1222.345030

	minor_axis_length	local_centroid_y	local_centroid_x	euler \
count	283.000000	283.000000	283.000000	283.000000
mean	80.548936	126.042499	48.882566	-0.667845
std	74.449697	138.949066	38.568380	3.493041
min	0.000000	0.000000	0.000000	-31.000000
25%	41.046784	33.224947	24.134395	-1.000000
50%	51.738385	52.634532	37.659745	1.000000
75%	103.064860	156.291373	62.515708	1.000000
max	471.031488	717.234813	162.353358	1.000000

	solidity
count	283.000000
mean	0.900393
std	0.142308
min	0.009825
25%	0.912838
50%	0.952869
75%	0.971785
max	1.000000

```
[ ]: import seaborn as sns

sns.pairplot(X)
```

1.2 The Models

```
[ ]: from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
```


```
[158]: # PCC

from collections import Counter
state_counts = Counter(Y)
df_state = pd.DataFrame.from_dict(state_counts, orient='index',
                                  columns=['class'])
df_state.plot(kind='bar', color='indigo', title='PCC')

num=(df_state['class']/df_state['class'].sum())**2
print("Population per class: {}".format(df_state))
print("1.25 * Proportion Chance Criterion: {}".format(1.25*100*num.sum()))
```

Population per class: class
A 58
B 50
C 57
D 63
E 55

1.25 * Proportion Chance Criterion: 25.13922011761915%



```
[159]: def train_leaves(X, Y):  
        """Run Grid Search to train the RF Classifier on different parameters."""  
        xtrain, xtest, ytrain, ytest = train_test_split(X, Y, test_size=0.25)  
  
        # Hyperparameters for tuning  
        params = {'n_estimators': [50, 100, 150, 200],  
                  'max_depth': range(2,5),  
                  'max_features': ['auto', 'sqrt', 'log2', 0.2, 0.3, 0.4]}  
  
        # Model training and evaluation  
        model = RandomForestClassifier()  
        model_cv = GridSearchCV(model, params)  
  
        model_cv.fit(xtrain, ytrain)  
  
        train_pred = model_cv.predict(xtrain)
```

```

test_pred = model_cv.predict(xtest)

test_score = model_cv.score(xtest, ytest)

#     print(model_cv.best_score_)
#     display(confusion_matrix(train_pred, ytrain))
#     display(confusion_matrix(test_pred, ytest))
return (model_cv.best_score_, test_score,
        confusion_matrix(train_pred, ytrain),
        confusion_matrix(test_pred, ytest))

```

1.2.1 Baseline

```

[173]: from tqdm import tqdm
cv_accuracy = []
test_accuracy = []
conf_train = []
conf_test = []

for i in tqdm(range(20)):
    score, tscore, ctrain, ctest = train_leaves(X, Y)
    #     print(f'Run {i}: {score}')
    cv_accuracy.append(score)
    test_accuracy.append(tscore)
    conf_train.append(ctrain)
    conf_test.append(ctest)

```

100%| | 20/20 [19:23<00:00, 58.17s/it]

```

[174]: print('Cross Validation Accuracy: ', np.mean(cv_accuracy))
print('Test Accuracy: ', np.mean(test_accuracy))

```

Cross Validation Accuracy: 0.8974584717607973
Test Accuracy: 0.8992957746478872

1.2.2 Perimeter-to-Area Ratio

```

[188]: X_pta = X.copy()

X_pta['pta_ratio'] = X['perimeter'] / X['area']

```

```

[189]: from tqdm import tqdm
cv_accuracy = []
test_accuracy = []
conf_train = []
conf_test = []

```

```

for i in tqdm(range(20)):
    score, tscore, ctrain, ctest = train_leaves(X_pta, Y)
    #     print(f'Run {i}: {score}')
    cv_accuracy.append(score)
    test_accuracy.append(tscore)
    conf_train.append(ctrain)
    conf_test.append(ctest)

```

100%| | 20/20 [19:50<00:00, 59.52s/it]

```

[190]: print('Cross Validation Accuracy: ', np.mean(cv_accuracy))
       print('Test Accuracy: ', np.mean(test_accuracy))

```

Cross Validation Accuracy: 0.8932668881506091
Test Accuracy: 0.8845070422535212

1.2.3 Major-Minor Axes Ratios

```

[191]: X_mma = X.copy()

       X_mma['mma_ratio'] = X['minor_axis_length'] / X['major_axis_length']

```

```

[192]: from tqdm import tqdm
       cv_accuracy = []
       test_accuracy = []
       conf_train = []
       conf_test = []

       for i in tqdm(range(20)):
           score, tscore, ctrain, ctest = train_leaves(X_mma, Y)
           #     print(f'Run {i}: {score}')
           cv_accuracy.append(score)
           test_accuracy.append(tscore)
           conf_train.append(ctrain)
           conf_test.append(ctest)

```

100%| | 20/20 [19:58<00:00, 59.91s/it]

```

[193]: print('Cross Validation Accuracy: ', np.mean(cv_accuracy))
       print('Test Accuracy: ', np.mean(test_accuracy))

```

Cross Validation Accuracy: 0.8975581395348838
Test Accuracy: 0.9056338028169014

```
[ ]:
```